

# Crop Disease Classification with CNN and NLP



## Project Overview

Our project uses Convolutional Neural Networks (CNNs) to automatically classify crop diseases from leaf images. By training on the PlantVillage dataset, the model learns to detect 14 types of diseases and healthy crops, helping improve diagnosis speed and accuracy in agriculture.

## Project Understanding

Crops are vulnerable to various diseases that can severely affect yields. Early detection is essential but not always accessible in rural areas. With CNNs, we can build an AI-powered tool that identifies diseases from simple leaf photos. The model is trained on thousands of labeled images and can achieve high accuracy, making it suitable for real-world use.

## Stakeholders

- **Farmers:** The primary users who benefit from fast and accurate disease diagnosis.
- **Agricultural Officers:** Can use the tool to support farmers in remote areas.
- **AgriTech Startups:** May integrate the model into mobile or web-based platforms.
- **Policy Makers:** Can use disease data trends to allocate resources effectively.
- **Students/Researchers:** Learn and improve AI models in the agricultural domain.

## Objectives

- Build a CNN model to classify crop diseases from leaf images.
- Develop NLP methods to analyze farmers' textual symptom descriptions.
- Integrate CNN and NLP outputs into a single multimodal diagnostic model.
- Compare multimodal performance to image-only and text-only models.

- Create a user-friendly API for farmers to upload photos and symptoms for instant diagnosis.

## Dataset

- Dataset: **PlantVillage**
- Classes:
  - Healthy
  - Rust
  - Blight
- Images size: 128x128 pixels
- Data split:
  - Training set: X images
  - Validation set: Y images
  - Test set: Z images

## Libraries Used

- TensorFlow / Keras
- NumPy
- Matplotlib
- Scikit-learn

## Data Loading & Preprocessing

```
In [13]: # Importing libraries
import os
import random
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras import layers
from tensorflow.keras.layers import (
    Conv2D, MaxPooling2D, Flatten, Dense, Dropout, BatchNormalization
)
from tensorflow.keras.utils import image_dataset_from_directory

# Set seeds for reproducibility
seed = 42
os.environ['PYTHONHASHSEED'] = str(seed)
```

```
tf.random.set_seed(seed)
np.random.seed(seed)
random.seed(seed)
```

## Data Loading

```
In [15]: # Loading the data
dataset_path = "PlantVillage"
data = tf.keras.utils.image_dataset_from_directory(
    dataset_path,
    image_size=(128, 128),
    batch_size=32,
    label_mode='categorical'
)
class_names = data.class_names
num_classes = len(class_names)
num_classes
```

Found 20638 files belonging to 15 classes.

Out[15]: 15

## Basic EDA (Exploratory Data Analysis) on dataset

```
In [17]: # Finding out how many imgs are in each folder
import os

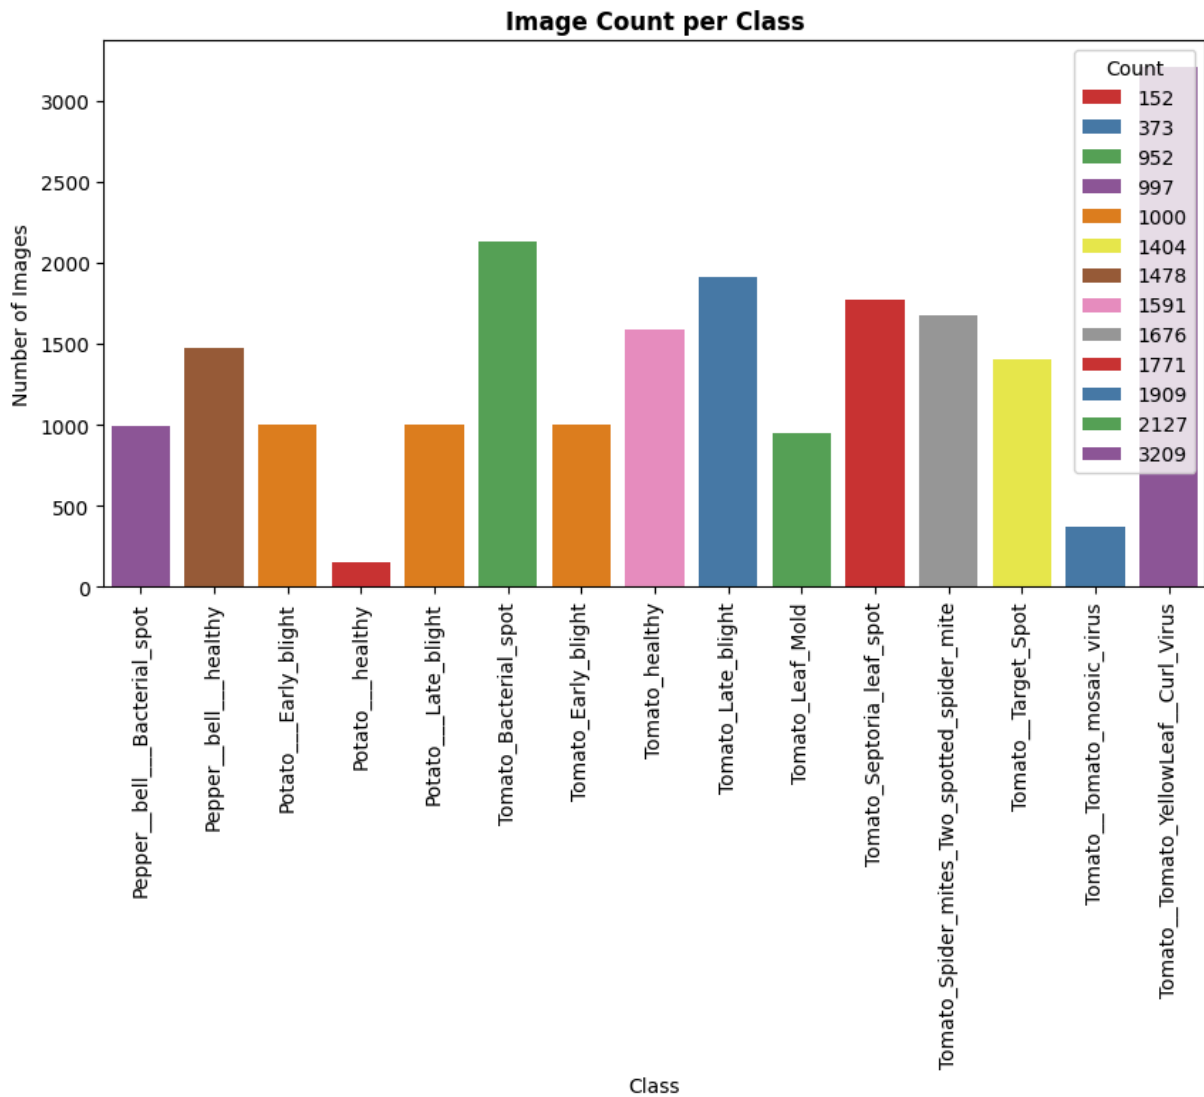
counts = {}
for folder in os.listdir(dataset_path):
    folder_path = os.path.join(dataset_path, folder)
    if os.path.isdir(folder_path):
        count = len(os.listdir(folder_path))
        counts[folder] = count

# Display results
for cls, count in counts.items():
    print(f"{cls}: {count} images")
```

Pepper\_\_bell\_\_Bacterial\_spot: 997 images  
Pepper\_\_bell\_\_healthy: 1478 images  
Potato\_\_Early\_blight: 1000 images  
Potato\_\_healthy: 152 images  
Potato\_\_Late\_blight: 1000 images  
Tomato\_Bacterial\_spot: 2127 images  
Tomato\_Early\_blight: 1000 images  
Tomato\_healthy: 1591 images  
Tomato\_Late\_blight: 1909 images  
Tomato\_Leaf\_Mold: 952 images  
Tomato\_Septoria\_leaf\_spot: 1771 images  
Tomato\_Spider\_mites\_Two\_spotted\_spider\_mite: 1676 images  
Tomato\_\_Target\_Spot: 1404 images  
Tomato\_\_Tomato\_mosaic\_virus: 373 images  
Tomato\_\_Tomato\_YellowLeaf\_\_Curl\_Virus: 3209 images

```
In [18]: # Plotting a distribution for the classes
counts_df = pd.DataFrame({
    "Class": list(counts.keys()),
    "Count": list(counts.values())
})

plt.figure(figsize=(10,5))
sns.barplot(data=counts_df, x="Class", y="Count", hue = 'Count', palette="Set1")
plt.xticks(rotation=90)
plt.title("Image Count per Class", weight = 'bold')
plt.ylabel("Number of Images")
plt.show()
```



## Visualize Sample Images

```
In [20]: for images, labels in data.take(1):
plt.figure(figsize=(10,10))
for i in range(9):
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(images[i].numpy().astype("uint8"))
    class_idx = np.argmax(labels[i])
    plt.title(class_names[class_idx])
    plt.tight_layout()
    plt.axis("off")
```

Tomato\_Septoria\_leaf\_spot



Tomato\_healthy



Tomato\_healthy



Tomato\_Leaf\_Mold



Tomato\_Leaf\_Mold



Tomato\_healthy



Tomato\_\_Tomato\_YellowLeaf\_Curl\_Virus



Tomato\_Late\_blight



Potato\_\_Early\_blight



```
In [21]: # Checking image size
from PIL import Image

image_sizes = []

for class_folder in os.listdir(dataset_path):
    class_path = os.path.join(dataset_path, class_folder)
    for img_file in os.listdir(class_path)[:10]: # limit to 10 per class for
        img_path = os.path.join(class_path, img_file)
        with Image.open(img_path) as img:
            image_sizes.append(img.size)

# Convert to DataFrame for EDA
sizes_df = pd.DataFrame(image_sizes, columns=["Width", "Height"])
print(sizes_df.describe())
```

	Width	Height
count	150.0	150.0
mean	256.0	256.0
std	0.0	0.0
min	256.0	256.0
25%	256.0	256.0
50%	256.0	256.0
75%	256.0	256.0
max	256.0	256.0

## Preprocessing

### Splitting to test and validations

In [24]: *# Trainset*

```
original_train_ds = tf.keras.utils.image_dataset_from_directory(
    dataset_path,
    validation_split=0.2,
    subset="training",
    seed=42,
    image_size=(128, 128),
    batch_size=32,
    label_mode='categorical'
)

# Validation set

val_ds = tf.keras.utils.image_dataset_from_directory(
    dataset_path,
    validation_split=0.2,
    subset="validation",
    seed=42,
    image_size=(128, 128),
    batch_size=32,
    label_mode='categorical'
)
```

Found 20638 files belonging to 15 classes.  
 Using 16511 files for training.  
 Found 20638 files belonging to 15 classes.  
 Using 4127 files for validation.

### Rescale Pixels

In [26]: *# Rescaling/Normalization*

```
normalization_layer = layers.Rescaling(1./255)

train_ds = original_train_ds.map(lambda x, y: (normalization_layer(x), y))
val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y))
```

### Data Augmentation

Data augmentation creates new images from your existing ones, helping reduce overfitting.

```
In [29]: data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.1),
    layers.RandomZoom(0.1),
])

# Apply augmentation during training
augmented_train_ds = train_ds.map(lambda x, y: (data_augmentation(x, training=True), y))
```

## Prefetching

Speed up our pipeline by prefetching our data

```
In [32]: AUTOTUNE = tf.data.AUTOTUNE

train_ds = train_ds.prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.prefetch(buffer_size=AUTOTUNE)
```

```
In [33]: # Checking our preprocessed data
for images, labels in augmented_train_ds.take(1):
    plt.figure(figsize=(10,10))
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy())
        plt.title("Augmented Image")
        plt.axis("off")
```



Augmented Image



Augmented Image



Augmented Image



Augmented Image



Augmented Image



Augmented Image



Augmented Image



Augmented Image



Augmented Image



```
In [34]: for images, labels in train_ds.take(1):  
         print("Images shape:", images.shape)  
         print("Labels shape:", labels.shape)
```

Images shape: (32, 128, 128, 3)  
Labels shape: (32, 15)

## Modeling

```
In [36]: model = Sequential([  
         layers.Input(shape=(128, 128, 3)),  
  
         layers.Conv2D(32, (3, 3), activation='relu'),  
         layers.MaxPooling2D((2, 2)),  
  
         layers.Conv2D(64, (3, 3), activation='relu'),
```

```

layers.MaxPooling2D((2, 2)),

layers.Conv2D(128, (3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)),

layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dropout(0.5),
layers.Dense(num_classes, activation='softmax')
])

```

```

In [37]: model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
model.summary()

```

**Model: "sequential\_1"**

Layer (type)	Output Shape	Par
conv2d (Conv2D)	(None, 126, 126, 32)	
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	
conv2d_1 (Conv2D)	(None, 61, 61, 64)	18
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	
conv2d_2 (Conv2D)	(None, 28, 28, 128)	73
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	
flatten (Flatten)	(None, 25088)	
dense (Dense)	(None, 128)	3,211
dropout (Dropout)	(None, 128)	
dense_1 (Dense)	(None, 15)	1

**Total params:** 3,306,575 (12.61 MB)











**Trainable params:** 3,306,575 (12.61 MB)

**Non-trainable params:** 0 (0.00 B)

```



In [38]: history = model.fit(train_ds, validation_data=val_ds, epochs=10)

```

Epoch 1/10  
**516/516**  **278s** 528ms/step - accuracy: 0.3482 - loss: 2.03  
 27 - val\_accuracy: 0.7274 - val\_loss: 0.8090  
 Epoch 2/10  
**516/516**  **225s** 436ms/step - accuracy: 0.6747 - loss: 1.00  
 62 - val\_accuracy: 0.8205 - val\_loss: 0.5250  
 Epoch 3/10  
**516/516**  **245s** 475ms/step - accuracy: 0.7207 - loss: 0.84  
 09 - val\_accuracy: 0.8379 - val\_loss: 0.4699  
 Epoch 4/10  
**516/516**  **239s** 464ms/step - accuracy: 0.7768 - loss: 0.67  
 30 - val\_accuracy: 0.8616 - val\_loss: 0.4143  
 Epoch 5/10  
**516/516**  **246s** 476ms/step - accuracy: 0.8092 - loss: 0.56  
 88 - val\_accuracy: 0.8788 - val\_loss: 0.3595  
 Epoch 6/10  
**516/516**  **243s** 470ms/step - accuracy: 0.8300 - loss: 0.50  
 21 - val\_accuracy: 0.8878 - val\_loss: 0.3217  
 Epoch 7/10  
**516/516**  **211s** 408ms/step - accuracy: 0.8552 - loss: 0.41  
 99 - val\_accuracy: 0.8883 - val\_loss: 0.3151  
 Epoch 8/10  
**516/516**  **228s** 441ms/step - accuracy: 0.8695 - loss: 0.36  
 64 - val\_accuracy: 0.9164 - val\_loss: 0.2570  
 Epoch 9/10  
**516/516**  **237s** 459ms/step - accuracy: 0.8906 - loss: 0.31  
 78 - val\_accuracy: 0.9193 - val\_loss: 0.2615  
 Epoch 10/10  
**516/516**  **215s** 417ms/step - accuracy: 0.8958 - loss: 0.29  
 91 - val\_accuracy: 0.9159 - val\_loss: 0.2565

## Model evaluation

```
In [40]: # Evaluate how the model performs across all validations
loss, accuracy = model.evaluate(train_ds)
print("Train Loss:", loss)
print("Train Accuracy:", accuracy)
print()
loss, accuracy = model.evaluate(val_ds)
print("Validation Loss:", loss)
print("Validation Accuracy:", accuracy)
```

**516/516**  **52s** 101ms/step - accuracy: 0.9755 - loss: 0.078  
 2  
 Train Loss: 0.06999358534812927  
 Train Accuracy: 0.9795893430709839  
  
**129/129**  **13s** 98ms/step - accuracy: 0.9167 - loss: 0.2431  
 Validation Loss: 0.2565278708934784  
 Validation Accuracy: 0.9159195423126221

### Training Performance:

- Train Accuracy: 96.3%
- Train Loss: 0.121

This means the model is doing very well on the training data — it correctly classifies most images and has a low error rate.

### Validation Performance:

- Validation Accuracy: 89.6%
- Validation Loss: 0.320

This shows the model is still performing well on unseen data, though the accuracy is slightly lower than training, which is expected.


## Plotting the model performance

```
In [43]: fig, axs = plt.subplots(1, 2, figsize=(16, 6))

# Accuracy
axs[0].plot(history.history['accuracy'], label='Train Accuracy')
axs[0].plot(history.history['val_accuracy'], label='Validation Accuracy')
axs[0].set_title('Accuracy Over Epochs', weight = 'bold')
axs[0].legend()

# Loss
axs[1].plot(history.history['loss'], label='Train Loss')
axs[1].plot(history.history['val_loss'], label='Validation Loss')
axs[1].set_title('Loss Over Epochs', weight = 'bold')
axs[1].legend()

plt.show()
```



## Confusion Matrix

```
In [45]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.metrics import ConfusionMatrixDisplay

# get class names
```

```

val_dir = "./PlantVillage"
class_names = sorted(
    entry.name for entry in os.scandir(val_dir) if entry.is_dir()
)

# get predictions
y_true = []
y_pred = []


for images, labels in val_ds:
    preds = model.predict(images, verbose=0)
    y_true.extend(np.argmax(labels.numpy(), axis=1))
    y_pred.extend(np.argmax(preds, axis=1))

y_true = np.array(y_true)
y_pred = np.array(y_pred)

# compute confusion matrix
cm = confusion_matrix(y_true, y_pred)

# plot
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
disp.plot(cmap='Greens', xticks_rotation=90)
plt.title("Confusion Matrix - CNN Model", weight = 'bold')
plt.show()

```



## Making Predictions

```
In [47]: # Take one batch from validation set
for images, labels in val_ds.take(1):
    predictions = model.predict(images)

    # Get predicted classes
    predicted_classes = np.argmax(predictions, axis=1)
    true_classes = np.argmax(labels.numpy(), axis=1)

    print("Predicted:", predicted_classes)
    print("Actual:", true_classes)
```

1/1 ————— 0s 120ms/step

Predicted: [11 14 6 6 5 12 10 1 10 10 10 1 3 10 9 12 1 1 8 10 3  
7 1 9

0 14 3 12 12 7 1 1]

Actual: [11 14 9 6 5 12 10 1 10 10 10 1 6 10 9 12 1 1 8 10 3 3  
1 9

0 14 3 12 12 6 1 1]

```
In [48]: # Get class names
class_names = original_train_ds.class_names

# Print class names instead of numeric labels
for images, labels in val_ds.take(1):
    predictions = model.predict(images)

    predicted_classes = np.argmax(predictions, axis=1)
    true_classes = np.argmax(labels.numpy(), axis=1)

    predicted_names = [class_names[i] for i in predicted_classes]
    true_names = [class_names[i] for i in true_classes]

    print("Predicted class names:", predicted_names)
    print()
    print("Actual class names:", true_names)
```

1/1 ————— 0s 135ms/step

Predicted class names: ['Tomato\_Early\_blight', 'Tomato\_Spider\_mites\_Two\_spotted\_spider\_mite', 'Tomato\_Septoria\_leaf\_spot', 'Pepper\_bell\_healthy', 'Tomato\_Bacterial\_spot', 'Pepper\_bell\_Bacterial\_spot', 'Tomato\_Target\_Spot', 'Potato\_Late\_blight', 'Tomato\_healthy', 'Tomato\_Late\_blight', 'Tomato\_Spider\_mites\_Two\_spotted\_spider\_mite', 'Tomato\_Bacterial\_spot', 'Tomato\_Spider\_mites\_Two\_spotted\_spider\_mite', 'Tomato\_Tomato\_YellowLeaf\_Curl\_Virus', 'Tomato\_Late\_blight', 'Tomato\_Bacterial\_spot', 'Tomato\_Bacterial\_spot', 'Tomato\_Septoria\_leaf\_spot', 'Tomato\_Early\_blight', 'Tomato\_Septoria\_leaf\_spot', 'Tomato\_Bacterial\_spot', 'Pepper\_bell\_healthy', 'Pepper\_bell\_Bacterial\_spot', 'Tomato\_Bacterial\_spot', 'Tomato\_Septoria\_leaf\_spot', 'Tomato\_Tomato\_YellowLeaf\_Curl\_Virus', 'Tomato\_Septoria\_leaf\_spot', 'Tomato\_Spider\_mites\_Two\_spotted\_spider\_mite', 'Tomato\_Early\_blight', 'Tomato\_Tomato\_YellowLeaf\_Curl\_Virus', 'Tomato\_healthy', 'Tomato\_Target\_Spot']

Actual class names: ['Tomato\_Early\_blight', 'Tomato\_Spider\_mites\_Two\_spotted\_spider\_mite', 'Tomato\_Early\_blight', 'Pepper\_bell\_healthy', 'Tomato\_Bacterial\_spot', 'Pepper\_bell\_Bacterial\_spot', 'Tomato\_Target\_Spot', 'Potato\_Late\_blight', 'Tomato\_healthy', 'Tomato\_Leaf\_Mold', 'Tomato\_Spider\_mites\_Two\_spotted\_spider\_mite', 'Tomato\_Bacterial\_spot', 'Tomato\_Spider\_mites\_Two\_spotted\_spider\_mite', 'Tomato\_Tomato\_YellowLeaf\_Curl\_Virus', 'Tomato\_Early\_blight', 'Tomato\_Bacterial\_spot', 'Tomato\_Bacterial\_spot', 'Tomato\_Septoria\_leaf\_spot', 'Tomato\_Septoria\_leaf\_spot', 'Tomato\_Septoria\_leaf\_spot', 'Tomato\_Bacterial\_spot', 'Pepper\_bell\_healthy', 'Pepper\_bell\_Bacterial\_spot', 'Tomato\_Bacterial\_spot', 'Tomato\_Septoria\_leaf\_spot', 'Tomato\_Tomato\_YellowLeaf\_Curl\_Virus', 'Tomato\_Target\_Spot', 'Tomato\_Spider\_mites\_Two\_spotted\_spider\_mite', 'Tomato\_Early\_blight', 'Tomato\_Tomato\_YellowLeaf\_Curl\_Virus', 'Tomato\_healthy', 'Tomato\_Target\_Spot']

## NATURAL LANGUAGE PROCESSING

For this project we are also considering farmers who might not have access cameras or might have a hard time navigating to where an image is located.

For this giving descriptions of how the crop looks like might help this kind of farmers by giving a description and our **NLP** model predicts if the crop disease

that might be affecting that crop

Since we did not have access to actual collected descriptions of the crops in our dataset we opted to synthesize 1751 records of actual descriptions from the internet which is in the ***crop\_descripton.ipynb*** file and saved it in ***synthetic\_data.csv*** file.

The process involved having a dictionary of descriptions for each class and yousing adjectives and prefixes so that the data did not have many duplicates. Out of 1751 records only 13 are duplicated.

## Loading, Cleaning and Data Exploration

```
In [53]: # Library Importations
import nltk
import re
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.preprocessing import LabelEncoder
```

```
In [54]: des_df = pd.read_csv('synthetic_data.csv')
des_df.head()
```



	description	crop	status	disease	recommended_pesticide
0	Agronomist notes: extensive Stems remain large...	Tomato	Unhealthy	Septoria_leaf_spot	Bacillus subtilis
1	Agronomist notes: notable Yellow patches appea...	Tomato	Unhealthy	Leaf_Mold	Biological predator
2	Farmer complains: rapid Overall appearance is ...	Pepper_bell	Healthy	healthy	Na
3	Farmer complains: notable Dark spots sometimes...	Potato	Unhealthy	Early_blight	Bacillus subtilis
4	Field observation shows: slight Plant stands t...	Tomato	Healthy	healthy	Na

In [55]: `des_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1750 entries, 0 to 1749
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   description                           1750 non-null   object
1   crop                                  1750 non-null   object
2   status                               1750 non-null   object
3   disease                              1750 non-null   object
4   recommended_pesticide                 1403 non-null   object
dtypes: object(5)
memory usage: 68.5+ KB
```

In [56]: `des_df.duplicated().sum()`

Out[56]: 16

In [57]: `des_duplicated_df = des_df[des_df.duplicated()]`  
`des_duplicated_df`

Out[57]:

	description	crop	status	disease	re
<b>319</b>	Agronomist notes: sudden Growth appears stunte...	Tomato	Unhealthy	Target_Spot	
<b>491</b>	Field observation shows: extensive Spots are m...	Tomato	Unhealthy	Target_Spot	
<b>682</b>	Farmer reports: clear Lesions spread quickly u...	Potato	Unhealthy	Late_blight	
<b>685</b>	Inspection reveals: sporadic Overall plant app...	Tomato	Unhealthy	Tomato_YellowLeaf_Curl_Virus	
<b>913</b>	Farmer reports: slight Large irregular brown p...	Tomato	Unhealthy	Late_blight	
<b>1098</b>	Farmer reports: rapid Growth becomes stunted a...	Tomato	Unhealthy	Tomato_YellowLeaf_Curl_Virus	
<b>1207</b>	According to the grower: visible Overall plant...	Tomato	Unhealthy	Tomato_YellowLeaf_Curl_Virus	
<b>1214</b>	According to the grower: moderate Yellow patch...	Tomato	Unhealthy	Leaf_Mold	
<b>1237</b>	Farmer reports: slight Edges of lesions appear...	Potato	Unhealthy	Late_blight	
<b>1311</b>	Inspection reveals: sudden Lower	Tomato	Unhealthy	Leaf_Mold	

	description	crop	status	disease	re
	leaves are af...				
1533	Agronomist notes: sudden Webbing may be visibl...	Tomato	Unhealthy	Spider_mites_Two_spotted_spider_mite	
1627	Inspection reveals: severe Plants appear free ...	Tomato	Healthy		healthy
1633	Field observation shows: slight Growth appears...	Pepper_bell	Healthy		healthy
1702	Farmer complains: sudden Older leaves are the ...	Potato	Unhealthy		Early_blight
1713	According to the grower: visible Virus spreads...	Tomato	Unhealthy	Tomato_mosaic_virus	
1725	Agronomist notes: notable Plants look generall...	Tomato	Unhealthy	Septoria_leaf_spot	

```
In [58]: des_df = des_df.drop_duplicates()
des_df.duplicated().sum()
```

```
Out[58]: 0
```

```
In [59]: des_df.shape
```

```
Out[59]: (1734, 5)
```

```
In [60]: des_df.isna().sum()
```

```
Out[60]: description      0
crop                    0
status                  0
disease                 0
recommended_pesticide    345
dtype: int64
```

```
In [61]: #Replacing None in our pesticide column with No pesticide Required
des_df['pesticide'] = des_df['recommended_pesticide']
des_df = des_df.drop(columns = ['recommended_pesticide'])

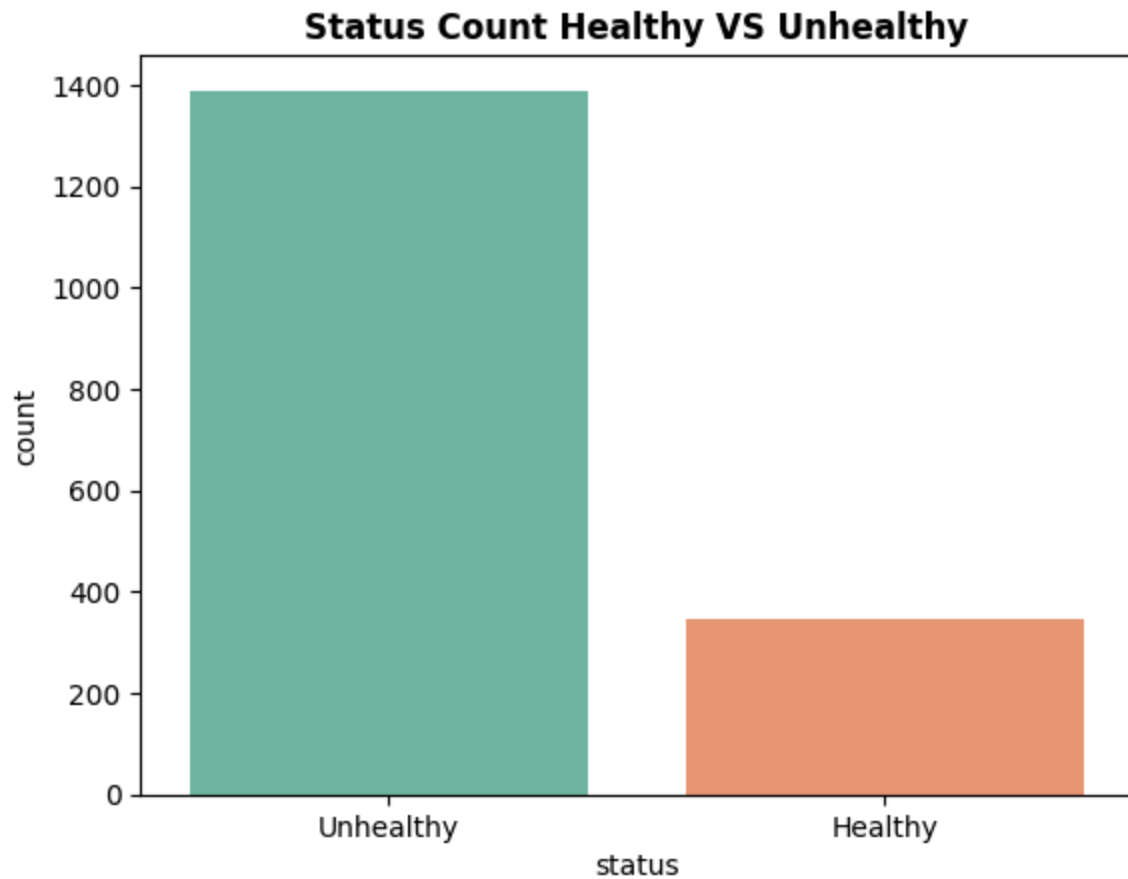
des_df = des_df.fillna('No pesticide needed')
des_df.head()
```

```
Out[61]:
```


	description	crop	status	disease	pesticide
0	Agronomist notes: extensive Stems remain large...	Tomato	Unhealthy	Septoria_leaf_spot	Bacillus subtilis
1	Agronomist notes: notable Yellow patches appea...	Tomato	Unhealthy	Leaf_Mold	Biological predators
2	Farmer complains: rapid Overall appearance is ...	Pepper_bell	Healthy	healthy	No pesticide needed
3	Farmer complains: notable Dark spots sometimes...	Potato	Unhealthy	Early_blight	Bacillus subtilis
4	Field observation shows: slight Plant stands t...	Tomato	Healthy	healthy	No pesticide needed

## Description Data Analysis

```
In [63]: # Count of healthy and Unhealthy status
sns.countplot(x = 'status', data = des_df, hue = 'status', palette = 'Set2')
plt.title('Status Count Healthy VS Unhealthy', weight = 'bold')
plt.show()
```

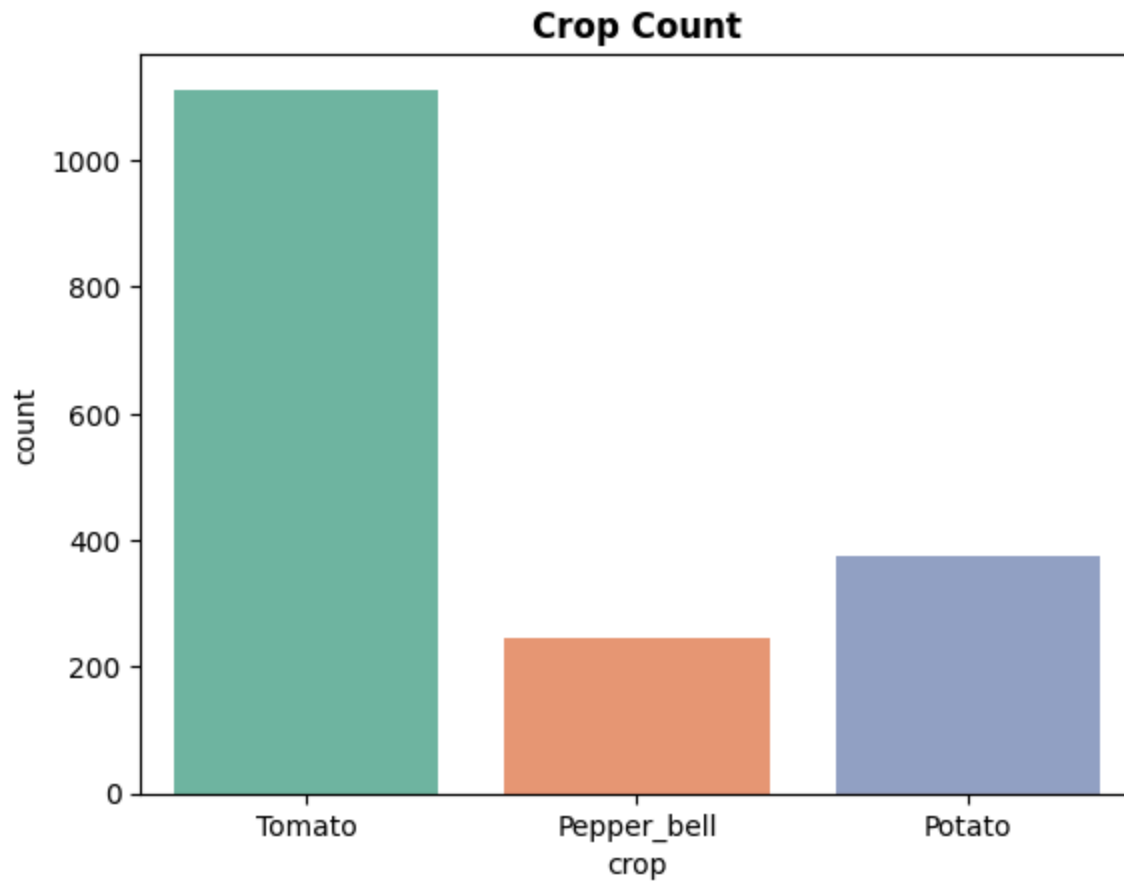


```
In [64]: # Disease Count
sns.countplot(y='disease', data=des_df, order=des_df['disease'].value_counts)
plt.title('Disease Class Distribution', weight = 'bold')
plt.show()
```

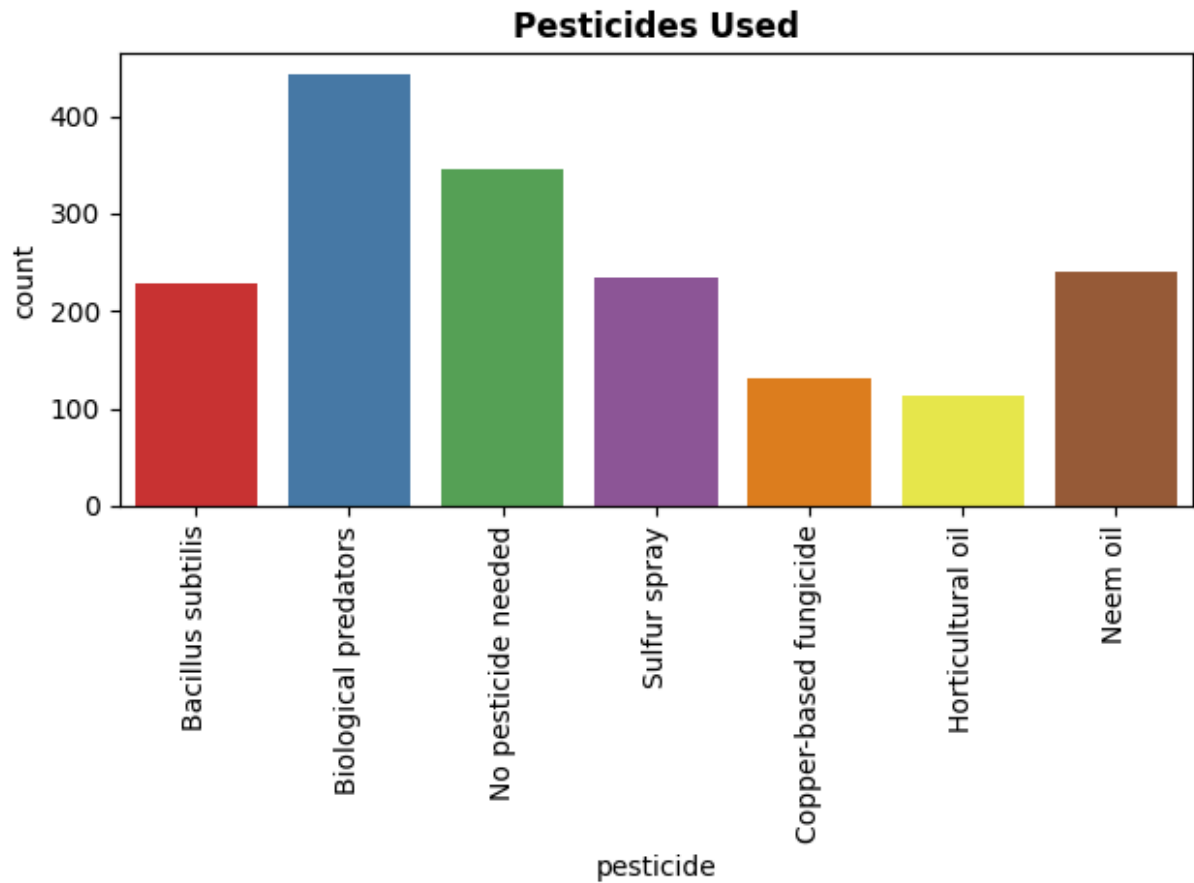


```
In [65]: # Crop count
sns.countplot(x = 'crop', data = des_df, hue = 'crop', palette = 'Set2')
```

```
plt.title('Crop Count', weight = 'bold')
plt.show()
```



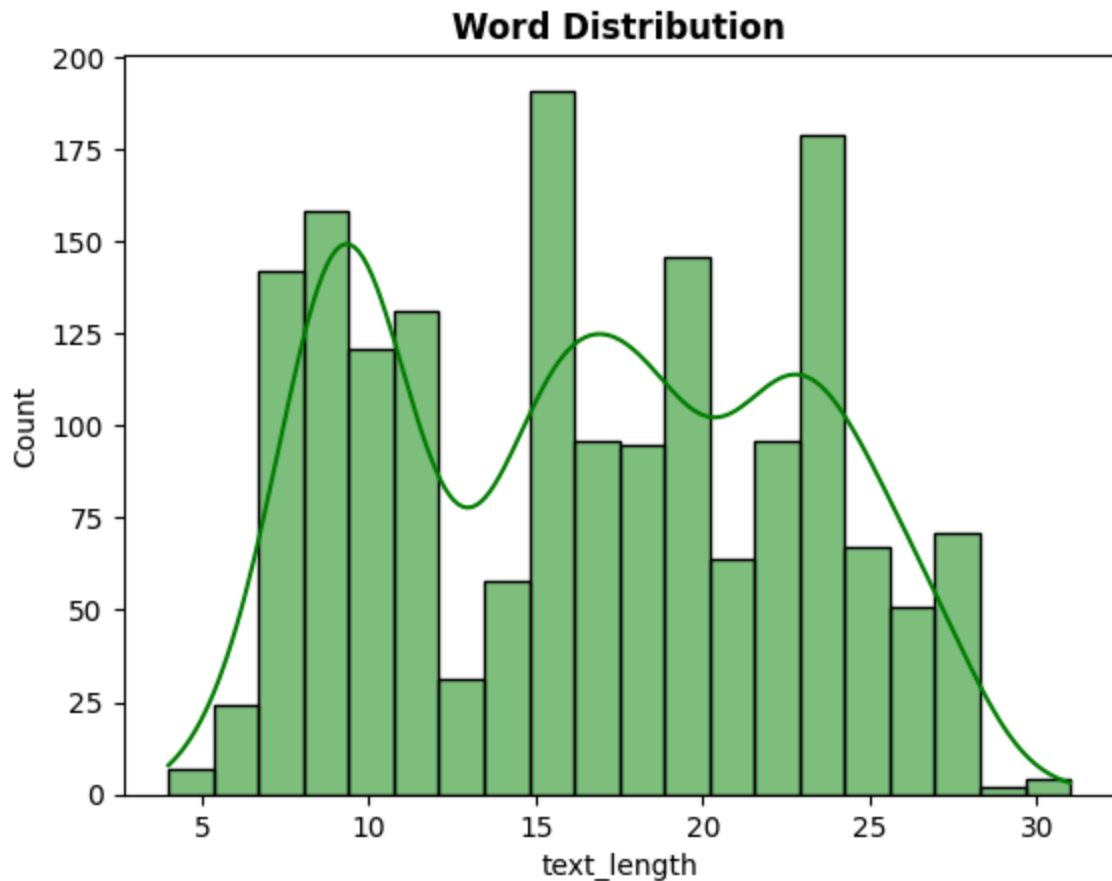
```
In [66]: # Pesticide Counts
sns.countplot(x = 'pesticide', data = des_df, hue = 'pesticide', palette = '
plt.title('Pesticides Used', weight = 'bold')
plt.xticks(rotation = 90)
plt.tight_layout()
plt.show()
```



## Text Analysis

```
In [68]: # Histogram of Words
des_df.loc[:, 'text_length'] = des_df['description'].apply(lambda x: len(x.s

sns.histplot(des_df['text_length'], bins = 20, kde = True, color = 'green')
plt.title('Word Distribution', weight = 'bold')
plt.show()
```



In [69]: `import spacy`

```
# Load tokenizer
nlp = spacy.load("en_core_web_sm")

text = "This is a test sentence."
tokens = [token.text for token in nlp(text)]
print(tokens)
```

`['This', 'is', 'a', 'test', 'sentence', '.']`

In [70]: `# Most Frequent Words`  
`from collections import Counter`

```
all_words = []

for desc in des_df['description']:
    doc = nlp(desc.lower())
    tokens = [token.text for token in doc if token.is_alpha and not token.is_punct]
    all_words.extend(tokens)

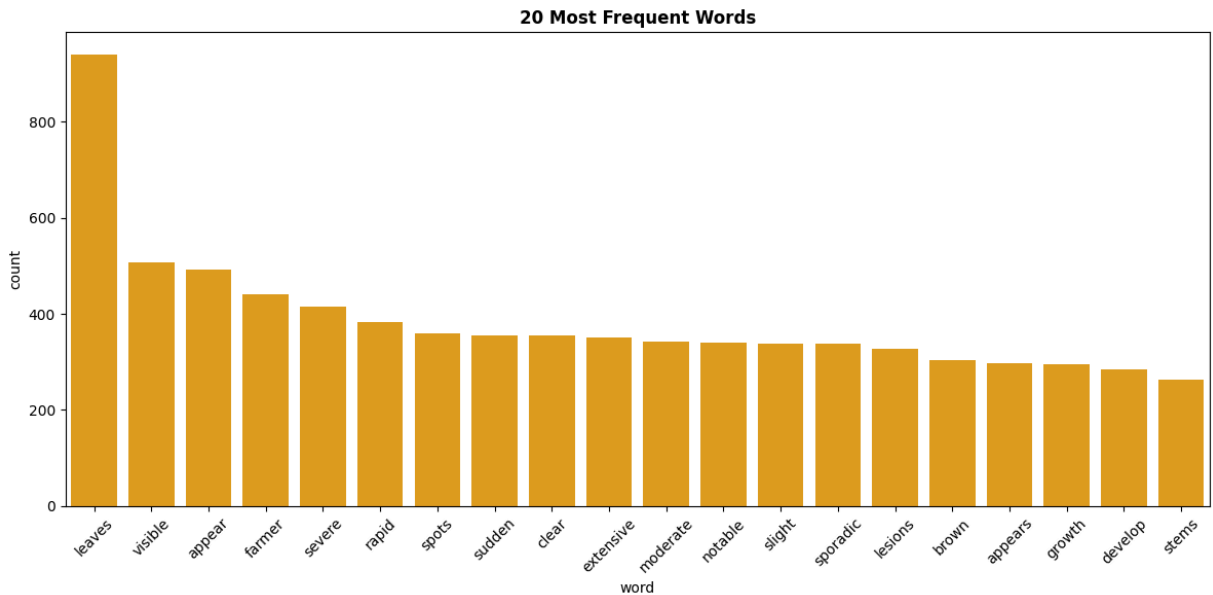
word_freq = Counter(all_words)
top_20_words = word_freq.most_common(20)

top_words_df = pd.DataFrame(top_20_words, columns=['word', 'count'])

# 5. Plotting the result
plt.figure(figsize=(12,6))
```



```
sns.barplot(data=top_words_df, x='word', y='count', color='orange')
plt.title('20 Most Frequent Words', weight='bold')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



## Feature Engineering

Since our target is the original 15 classes we need to make a new column with those classe with everything falling to where it is supposed to be.

For this we will need to combine our crop and disease column

```
In [73]: des_df['target_class'] = des_df['crop'] + '___' + des_df['disease']
des_df['target_class'].head()
```

```
Out[73]: 0    Tomato___Septoria_leaf_spot
1         Tomato___Leaf_Mold
2    Pepper_bell___healthy
3    Potato___Early_blight
4         Tomato___healthy
Name: target_class, dtype: object
```

```
In [74]: des_df.head()
```

	description	crop	status	disease	pesticide	text_length
0	Agronomist notes: extensive Stems remain large...	Tomato	Unhealthy	Septoria_leaf_spot	Bacillus subtilis	20
1	Agronomist notes: notable Yellow patches appea...	Tomato	Unhealthy	Leaf_Mold	Biological predators	10
2	Farmer complains: rapid Overall appearance is ...	Pepper_bell	Healthy	healthy	No pesticide needed	24
3	Farmer complains: notable Dark spots sometimes...	Potato	Unhealthy	Early_blight	Bacillus subtilis	9
4	Field observation shows: slight Plant stands t...	Tomato	Healthy	healthy	No pesticide needed	21

```
In [75]: print(f'Null Count: {des_df.isna().sum().sum()}')
print(f'Duplicate Count: {des_df.duplicated().sum()}')
```

Null Count: 0  
Duplicate Count: 0

## Text Preprocessing

```
In [77]: nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\PINCHEZZ\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\PINCHEZZ\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\PINCHEZZ\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

Out[77]: True

```
In [78]: stop_words = set(stopwords.words('english'))
         lemmatizer = WordNetLemmatizer()

         # Text cleaner function
         def clean_text_spacy(text):
             # Lowercase and remove non-letter characters
             text = re.sub(r'^a-zA-Z\s', '', text.lower())

             # Process the text
             doc = nlp(text)

             # Lemmatize, remove stop words and non-alphabetic tokens
             tokens = [token.lemma_ for token in doc if token.is_alpha and not token.

             return " ".join(tokens)
```

```
In [79]: # Applying our function to our DF
         des_df['clean_description'] = des_df['description'].apply(clean_text_spacy)
```

```
In [80]: des_df.head()
```

Out[80]:

	description	crop	status	disease	pesticide	text_length
--	-------------	------	--------	---------	-----------	-------------

0	Agronomist notes: extensive Stems remain large...	Tomato	Unhealthy	Septoria_leaf_spot	Bacillus subtilis	20
---	--	--------	-----------	--------------------	----------------------	----

1	Agronomist notes: notable Yellow patches appea...	Tomato	Unhealthy	Leaf_Mold	Biological predators	10
---	--	--------	-----------	-----------	-------------------------	----

2	Farmer complains: rapid Overall appearance is ...	Pepper_bell	Healthy	healthy	No pesticide needed	24
---	--	-------------	---------	---------	---------------------------	----

3	Farmer complains: notable Dark spots sometimes...	Potato	Unhealthy	Early_blight	Bacillus subtilis	9
---	---	--------	-----------	--------------	----------------------	---

4	Field observation shows: slight Plant stands t...	Tomato	Healthy	healthy	No pesticide needed	21
---	---	--------	---------	---------	---------------------------	----

```
In [81]: # Repositioning our target class to be last in our data frame
des_df = des_df[[col for col in des_df.columns if col != 'target_class']] + [

#Checking
des_df.head()
```

```
Out[81]:
```

	description	crop	status	disease	pesticide	text_length
0	Agronomist notes: extensive Stems remain large...	Tomato	Unhealthy	Septoria_leaf_spot	Bacillus subtilis	20
1	Agronomist notes: notable Yellow patches appea...	Tomato	Unhealthy	Leaf_Mold	Biological predators	10
2	Farmer complains: rapid Overall appearance is ...	Pepper_bell	Healthy	healthy	No pesticide needed	24
3	Farmer complains: notable Dark spots sometimes...	Potato	Unhealthy	Early_blight	Bacillus subtilis	9
4	Field observation shows: slight Plant stands t...	Tomato	Healthy	healthy	No pesticide needed	21

## TF-IDF Vectorization

```
In [83]: # Initialize vectorizer
tfidf = TfidfVectorizer(max_features=2000, ngram_range=(1,2))

#Fit with clean_description
x_tfidf = tfidf.fit_transform(des_df['clean_description'])

x_tfidf.shape
```

```
Out[83]: (1734, 2000)
```

## Label Encoding

```
In [85]: le = LabelEncoder()
y_encoded = le.fit_transform(des_df['target_class'])
```

```
# Check for correct mapping
class_mapping = dict(zip(le.classes_, le.transform(le.classes_)))
print(class_mapping)
```

```
{'Pepper_bell__Bacterial_spot': 0, 'Pepper_bell__healthy': 1, 'Potato__Early_blight': 2, 'Potato__Late_blight': 3, 'Potato__healthy': 4, 'Tomato__Bacterial_spot': 5, 'Tomato__Early_blight': 6, 'Tomato__Late_blight': 7, 'Tomato__Leaf_Mold': 8, 'Tomato__Septoria_leaf_spot': 9, 'Tomato__Spider_mites_Two_spotted_spider_mite': 10, 'Tomato__Target_Spot': 11, 'Tomato__Tomato_YellowLeaf_Curl_Virus': 12, 'Tomato__Tomato_mosaic_virus': 13, 'Tomato__healthy': 14}
```

## Train/Test Split

```
In [87]: # Splitting our data into train test and split for our model
X_train, X_test, y_train, y_test = train_test_split(x_tfidf, y_encoded, test
```

## Modeling

### Logistic Base Model


```
In [90]: # Build and fit the model
model_0 = LogisticRegression(max_iter=1000)
model_0.fit(X_train, y_train)

# Make predictions
y_preds = model_0.predict(X_test)
y_probs = model_0.predict_proba(X_test)
```

```
In [91]: # Evaluate model
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.preprocessing import label_binarize
from sklearn.metrics import roc_curve, auc

cm = confusion_matrix(y_test, y_preds)
disp = ConfusionMatrixDisplay(confusion_matrix = cm, display_labels=le.classes_)
disp.plot(cmap = 'Greens')
plt.title('Confusion Matrix For Base Model', weight = 'bold')
plt.xticks(rotation = 90)
plt.show()

print(classification_report(y_test, y_preds, target_names=le.classes_))
acc = accuracy_score(y_test, y_preds)
print(f"\nAccuracy: {acc:.4f}")
```



support		precision	recall	f1-score
26	Pepper_bell___Bacterial_spot	1.00	1.00	1.00
23	Pepper_bell___healthy	1.00	1.00	1.00
24	Potato___Early_blight	1.00	0.92	0.96
26	Potato___Late_blight	1.00	1.00	1.00
25	Potato___healthy	1.00	1.00	1.00
23	Tomato___Bacterial_spot	1.00	1.00	1.00
20	Tomato___Early_blight	0.91	1.00	0.95
25	Tomato___Late_blight	1.00	1.00	1.00
24	Tomato___Leaf_Mold	1.00	1.00	1.00
22	Tomato___Septoria_leaf_spot	1.00	1.00	1.00
23	Tomato___Spider_mites_Two_spotted_spider_mite	1.00	1.00	1.00
20	Tomato___Target_Spot	1.00	1.00	1.00
23	Tomato___Tomato_YellowLeaf_Curl_Virus	1.00	1.00	1.00
22	Tomato___Tomato_mosaic_virus	1.00	1.00	1.00
21	Tomato___healthy	1.00	1.00	1.00
347	accuracy			0.99
347	macro avg	0.99	0.99	0.99
347	weighted avg	0.99	0.99	0.99

Accuracy: 0.9942

```
In [92]: y_test_bin = label_binarize(y_test, classes=list(range(len(le.classes_)))) #

fpr = {}
tpr = {}
roc_auc = {}

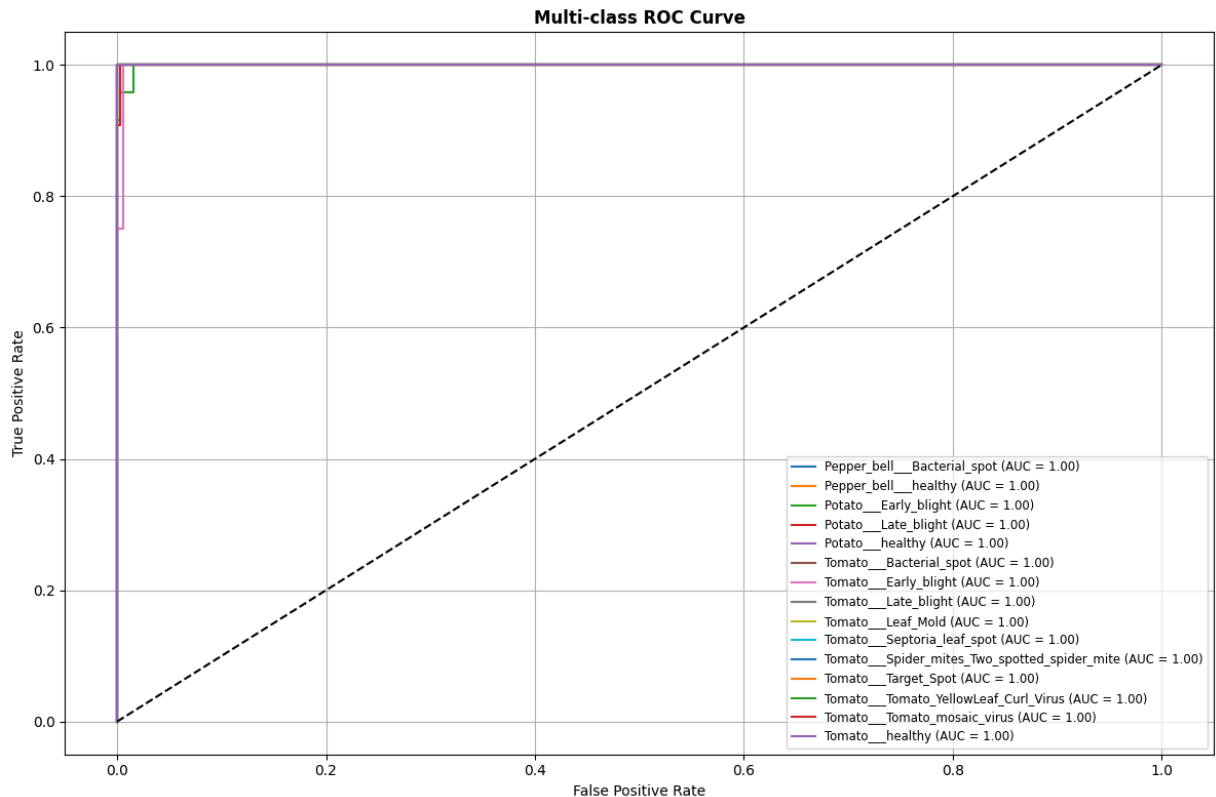
# Loop through each class
for i in range(len(le.classes_)):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_probs[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot
```

```
plt.figure(figsize=(12, 8))

for i in range(len(le.classes_)):
    plt.plot(fpr[i], tpr[i], label=f'{le.classes_[i]} (AUC = {roc_auc[i]:.2f})')

plt.plot([0, 1], [0, 1], 'k--') # diagonal line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Multi-class ROC Curve', weight='bold')
plt.legend(loc='lower right', fontsize='small')
plt.grid(True)
plt.tight_layout()
plt.show()
```



## Multinomial Naive Bayes Model

```
In [94]: #Library importation
from sklearn.pipeline import Pipeline
from sklearn.naive_bayes import MultinomialNB
from sklearn.base import BaseEstimator, TransformerMixin
import spacy
import joblib

nlp = spacy.load('en_core_web_sm')
```

```
In [95]: # Train_Test_split our data
X_train, X_test, y_train, y_test = train_test_split(
    des_df["clean_description"],
    des_df["target_class"],
    test_size=0.2,
```




```
    random_state=42,  
    stratify=des_df["target_class"]  
)
```

```
In [96]: # Initialize the pipeline  
nb_pipeline = Pipeline([  
    ('tfidf', TfidfVectorizer(  
        max_features=2000,  
        ngram_range=(1, 2)  
    )),  
    ('clf', MultinomialNB())  
])
```

```
In [97]: # Fit the pipeline on training data  
nb_pipeline.fit(X_train, y_train)
```

Out[97]:



## Predictions and Evaluations

```
In [99]: # Predict on the test set  
y_pred = nb_pipeline.predict(X_test)  
  
# Evaluate  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print("\nClassification Report:\n")  
print(classification_report(y_test, y_pred))
```


Accuracy: 0.9942363112391931

Classification Report:

support		precision	recall	f1-score
26	Pepper_bell___Bacterial_spot	1.00	1.00	1.00
23	Pepper_bell___healthy	1.00	1.00	1.00
24	Potato___Early_blight	1.00	0.92	0.96
26	Potato___Late_blight	1.00	1.00	1.00
25	Potato___healthy	1.00	1.00	1.00
23	Tomato___Bacterial_spot	1.00	1.00	1.00
20	Tomato___Early_blight	0.91	1.00	0.95
25	Tomato___Late_blight	1.00	1.00	1.00
24	Tomato___Leaf_Mold	1.00	1.00	1.00
22	Tomato___Septoria_leaf_spot	1.00	1.00	1.00
Tomato___Spider_mites_Two_spotted_spider_mite		1.00	1.00	1.00
23	Tomato___Target_Spot	1.00	1.00	1.00
20	Tomato___Tomato_YellowLeaf_Curl_Virus	1.00	1.00	1.00
23	Tomato___Tomato_mosaic_virus	1.00	1.00	1.00
22	Tomato___healthy	1.00	1.00	1.00
21				
347	accuracy			0.99
347	macro avg	0.99	0.99	0.99
347	weighted avg	0.99	0.99	0.99

Plot Confusion Matrix

```
In [101]: cm = confusion_matrix(y_test, y_pred, labels=nb_pipeline.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=nb_pipeline.classes_)
disp.plot(xticks_rotation=90, cmap='Greens')
plt.title('Confusion Matrix for Multinomial Naive Bayes', weight='bold')
plt.show()
```



## Saving Models and Usable Variables

```
In [103... # Saving CNN Model
model.save("cnn_model.keras")

#Saving NLP Model
joblib.dump(nb_pipeline, 'mnb_nlp_pipeline.pkl')
```

```
Out[103... ['mnb_nlp_pipeline.pkl']
```

```
In [ ]:
```